

Parallel Computing, 2022S

Assignment 2

Solving Heat Equation in 2D


The heat equation is a Partial Differential Equation

$$\frac{\partial u}{\partial t} = \alpha \left(\frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} \right)$$


Jacobi iterative method:

$$v_{m,l}^{n+1} = \frac{1}{4} (v_{m+1,l}^n + v_{m-1,l}^n + v_{m,l+1}^n + v_{m,l-1}^n) - \frac{h^2}{4} f_{ml} \quad \text{Zero}$$

- Working NxM matrix
- Five-point stencil
- Calculation of averages



Example output



Five-point stencil

Your Task

Realize an efficient parallel implementation of the serial code for iteratively solving heat equation using MPI.

1. **Distribute data over all processes**

- Vertical, horizontal, 2d?
- Also initialize in a distributed fashion
- **Each process gets only a part of the matrix**
- Support $N \neq M$
 - e.g., have larger last block

2. **Use point-to-point communication to communicate the overlapping regions**

3. **Make sure that you get the termination criteria correctly**



- `diffnorm` has to be the same as if executed with the sequential code

4. **Verify/compare to the sequential version**

- Use **collectives** to transfer data to rank 0, and then compare!
 - You can use the `verify` function from `a2-helpers.hpp`
- **Results must be the same!**

Hints: First support $N=M$ and then improve

Useful MPI routines: `MPI_Send/Recv`, `MPI_Allreduce`, `MPI_Gatherv`



Your Task

Use MPI to develop an efficient parallel code that can utilize all available nodes on a cluster

Develop anywhere, test on ALMA

1. Use up to all nodes and all cores
 2. Allow different sizes of matrix and also support $M \neq N$
 - But you can assume that matrix is larger than some minimum
 3. Try the following code parameters:
 - a. `mpirun ./a2 --m 2688 --n 4096 --epsilon 0.001 --max-iterations 1000`
 - b. `mpirun ./a2 --m 2688 --n 4096 --epsilon 0.001 --max-iterations 2000`
 - c. `mpirun ./a2 --m 1152 --n 1152 --epsilon 0.001 --max-iterations 1000`
 4. Try the following cluster configurations:
 - a. 1 node, 16 processes
 - b. 2 nodes, 32 processes
 - c. 4 nodes, 64 processes
 - d. 6 nodes, 96 processes
- Achieve a speedup of 12+ using all available nodes on the Alma system (using 3.a. and 4.d) and observe the speedup when using other configurations.
 - Try to explain the differences (if any) in the performance as M, N, and max-iterations change

Note that in most cases the code will reach the maximum number of iterations in the current setup

Your Task

Write a **report** and include:

1. Speedup results (**graphs** + **tables**)
 - Measure performance (computation and gather separately – as pointed out in the sequential code)
 - You can consider avoiding measuring MPI_Init() time
 - Use MPI_Wtime() for time measurements
 - Compare speedups for different matrix sizes (given in the last slide)
2. Source codes + Slurm job scripts
3. Discussion
 - Performance
 - Communication
 - and effects of the communication on the performance
 - Bottlenecks?
 - Possible improvements?

Sequential Code, Template, Compilation

Given in Moodle

- Main part of the source code: [a2.cpp](#)
- Helper functions and data structures: [a2-helpers.cpp](#)
- Compile (Alma):
 - `g++ -O2 -o a2 a2.cpp (sequential version)`
 - `mpic++ -O2 -o a2 a2.cpp (MPI)`
- Run (Usage):
 - Sequential
 - `./a2 --n <columns> --m <rows> [--max-iterations <int> --epsilon <double>]`
 - To run your MPI program
 - `mpirun -np ./a2 --n <columns> --m <rows> [--max-iterations <int> --epsilon <double>]`
 - Use Slurm to run on Alma

Sequential Code Notes (a2-helpers.hpp)

Contains a custom data structure “Mat”

- Contiguous 2d vector
- Note the “(“ and “)” instead of “[“ and “]”

And a set of helper functions that you do **not** need to modify:

- `void process_input(int argc, char **argv, int& N, int& M, int& max_iterations, double& epsilon);`
 - Process input arguments
- `void save_to_disk(Mat& U, std::string filename);`
 - Save data to disk (so it can be converted to an image)
- `void print_mat(Mat& m);`
 - Prints to std output
- `bool verify(Mat& a, Mat& b);`
 - Compares two give matrices
- `void heat2d_sequential(int max_iterations, double epsilon, Mat& U, int& iteration_count);`
 - Sequential version of the code (for verification)

```
Mat mat(height, width);
for (int i=0; i<mat.height; i++) {
    for (int j = 0; j < mat.width; j++) {
        mat(i,j) = f(x);
    }
}
```

Slurm

Job scheduling system for clusters

<https://slurm.schedmd.com/>

→ You submit a job script which tells the system how to execute your MPI program on the cluster

sbatch

- Submit a batch script to Slurm

squeue

- view information about jobs located in the Slurm scheduling queue

sstat

- Display various status information of a running job/step

sinfo (-N)

- View information about Slurm nodes and partitions

scancel

- Used to signal jobs or job steps that are under the control of Slurm (cancel a job with some ID)

Job Script

```
#!/bin/bash
```

```
#SBATCH -N 4
```

← 4 nodes

```
#SBATCH --ntasks 16
```

← MPI tasks

} 16 processes over 4 nodes

```
mpirun --report-bindings ./a2
```

↑ Optional -> lets you know where processes were executed

The result will be saved in slurm-6235.out when complete

```
> sbatch jobscript.sh
Submitted batch job 6235
```

```
> squeue
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
6236	all	jobscrip	johndoe	R	0:11	4	alma[01-04]

Alma System

Frontend + 6 worker nodes

Node:

- 2x Intel Xeon octa-core X5550 2.66Ghz
- 2x nVidia Tesla K20m
- 128GB RAM
- CentOS

Access via Frontend nodes

Shared file system

**Submit jobs scheduled
to run on compute nodes**

